

CSE 344 Computer Vision Assignment 2

Name : Arka Sarkar
Roll Number : 2018222

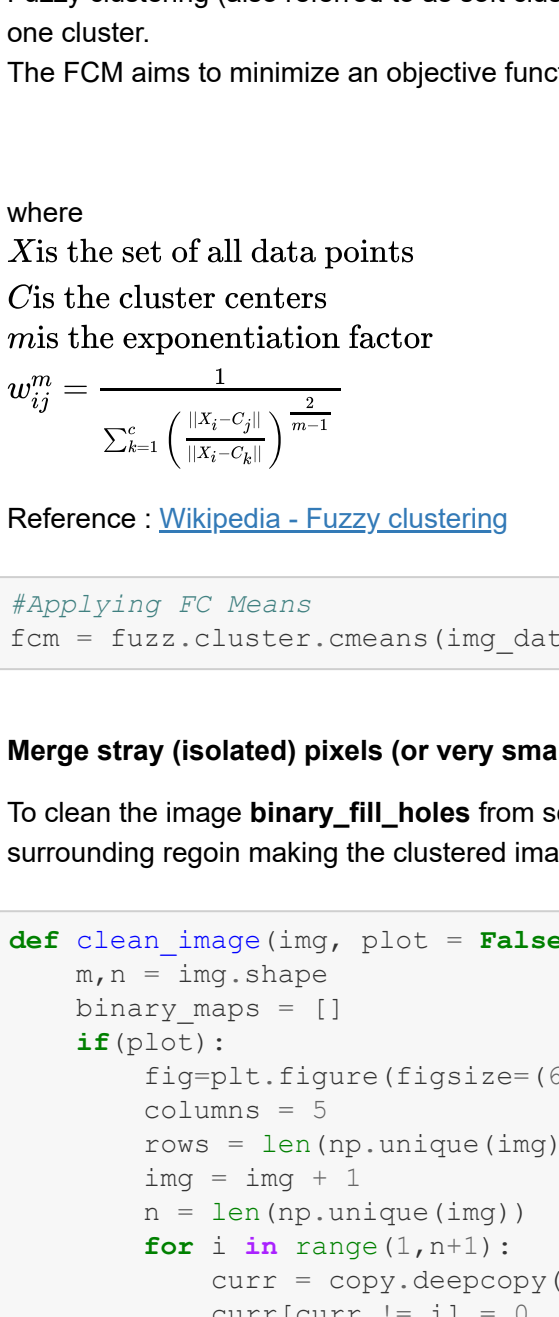
In [1]: `#dependencies Required`

```
import numpy as np
import cv2
import matplotlib.pyplot as plt
from IPython.display import display
from PIL import Image
import skimage as ssk
from sklearn.metrics import silhouette_score
from tqdm import tqdm
from skimage.segmentation import slic
from skimage.measure import label_boundaries
from skimage.util import img_as_float
from skimage import io
from scipy.spatial.distance import cdist
import scipy.stats
import copy
import math
from scipy.integrate import quad
from sklearn.cluster import KMeans
from scipy import ndimage
```

Question 1

Write a program to implement a region segmentation algorithm using the fuzzy c-means algorithm on normalized 'RGBY' data of an image. Merge stray (isolated) pixels (or very-small regions) to their surrounding regions. [3 marks]

In [2]: `img = cv2.imread('cap1.PNG')
fig=plt.figure(figsize=(6, 6))
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.show()`



In [3]: `#Making the Dataset`

```
img = np.array(img)
img_data = np.zeros((img.shape[0], img.shape[1], 5))
for i in range(img.shape[0]):
    for j in range(img.shape[1]):
        img_data[i,j,0:3] = img[i,j]/255
        img_data[i,j,3] = 1/img.shape[0]
        img_data[i,j,4] = 1/img.shape[1]
img_data = img_data.reshape(((img.shape[0]*img.shape[1], 5)))
```

Fuzzy C Means Clustering

Fuzzy clustering (also referred to as soft clustering or soft k-means) is a form of clustering in which each data point can belong to more than one cluster.

The FCM aims to minimize an objective function:

$$\underset{c}{\operatorname{argmin}} \sum_{i=1}^n \sum_{j=1}^c w_{ij}^m \times \|X_i - C_j\|^2$$

where

X_i is the set of all data points

C_j is the cluster centers

w_{ij} is the exponentiation factor

$$w_{ij}^m = \frac{1}{\sum_{k=1}^c \left(\frac{\|X_i - C_j\|}{\|X_i - C_k\|} \right)^{\frac{1}{1-m}}}$$

Reference : [Wikipedia - Fuzzy clustering](#)

In [4]: `#Applying FC Means`
`fcm = fuzz.cluster.cmeans(img_data, 25, 2, error=0.05, maxiter=1000, init=None)`

Merge stray (isolated) pixels (or very small regions) to their surrounding regions.

To clean the image binary_fill_holes from skimage image was used on every binary cluster iteratively to fill small pixels regions to the surrounding region making the clustered image more smoother.

In [19]: `def clean_image(img, plot = False):
 n,n = img.shape
 binary_maps = []
 if(plot):
 fig=plt.figure(figsize=(60, 60))
 columns = 5
 rows = len(np.unique(img))/5
 for i in range(binary_maps):
 n = len(np.uniqe(img))
 for i in range(1,n+1):
 curr = copy.deepcopy(img)
 curr[curr == i] = 0
 curr[curr == i] = 1
 filled = ndimage.binary_fill_holes(curr).astype(int)
 idx = np.where(filled == 1)
 img[idx[0],idx[1]] = i
 fig.add_subplot(rows, columns, i)
 plt.imshow(filled,astype(int))
 fig.tight_layout()
 plt.show()
 return img-1
 img = img + 1
 n = len(np.unique(img))
 for i in range(1,n+1):
 curr = copy.deepcopy(img)
 curr[curr == i] = 0
 curr[curr == i] = 1
 filled = ndimage.binary_fill_holes(curr).astype(int)
 idx = np.where(filled == 1)
 img[idx[0],idx[1]] = i
 return img-1`

In [6]: `cluster_centers = fcm[0]
prob_matrix = np.argmax(prob_matrix, axis = 0)
pred_matrix = cluster_centers[pred_matrix]
clustered_image = pred_matrix[:,0:3]
clustered_image = clustered_image.reshape(((img.shape[0],img.shape[1], 3)))
clustered_image = clustered_image.astype('float32')
pred_matrix = np.argmax(prob_matrix, axis = 0)
cleaned_img = clean_image(np.squeeze(pred_matrix).reshape(img.shape[0],img.shape[1]))
cleaned_plus_img = cluster_centers[cleaned_img[:,0:3]]
cleaned_plus_img = cleaned_plus_img.reshape(((img.shape[0],img.shape[1], 3)))
cleaned_plus_img = cleaned_plus_img.astype('float32')
fig=plt.figure(figsize=(8, 8))
columns = 2
rows = 1
fig.add_subplot(rows, columns, 1)
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
plt.imshow(cv2.cvtColor(clustered_image, cv2.COLOR_BGR2RGB))
plt.title("Un-Cleaned Clustered Image")
fig.add_subplot(rows, columns, 2)
plt.imshow(cv2.cvtColor(cleaned_plus_img, cv2.COLOR_BGR2RGB))
plt.title("Cleaned Clustered Image")
fig.tight_layout()
plt.show()`

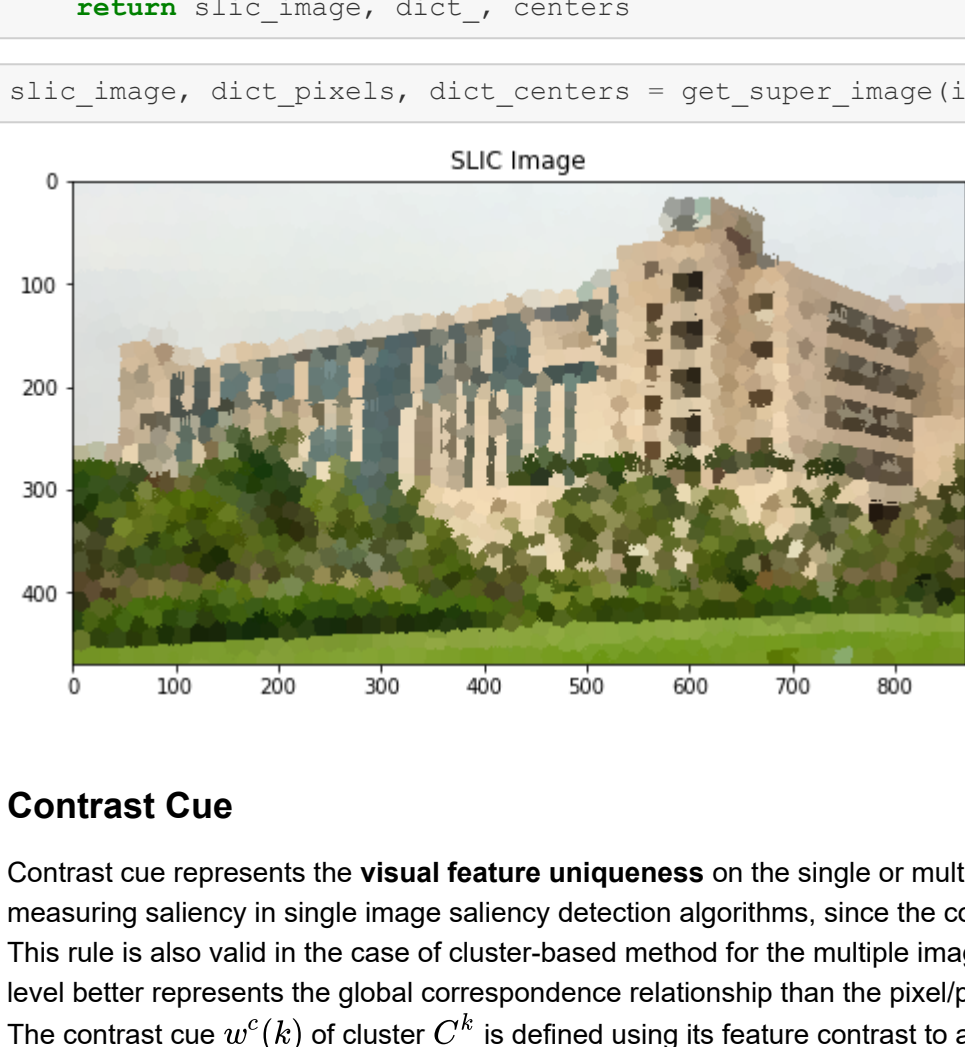


Question 2

Write a program to obtain the spatial and contrast cues using SLIC superpixels of an image instead of pixels. [3 marks]

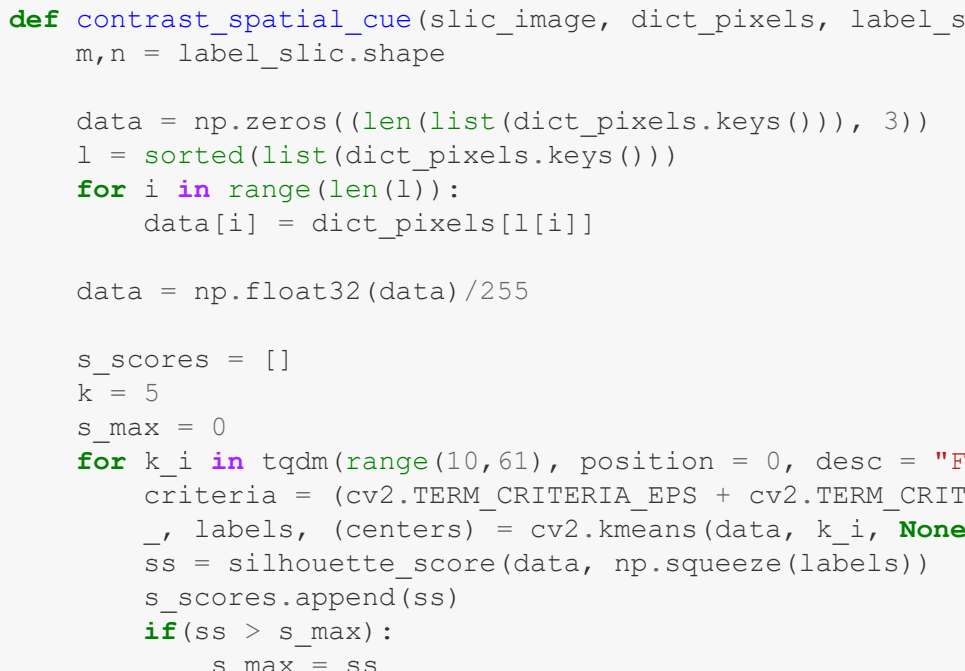
In [7]: `img_path = '15558014721_E7jYws_iit_d.jpg'
image = cv2.imread(img_path)
plt.figure(figsize = (8,8))
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))`

Out[7]: `<matplotlib.image.AxesImage at 0x24b589ddc8>`



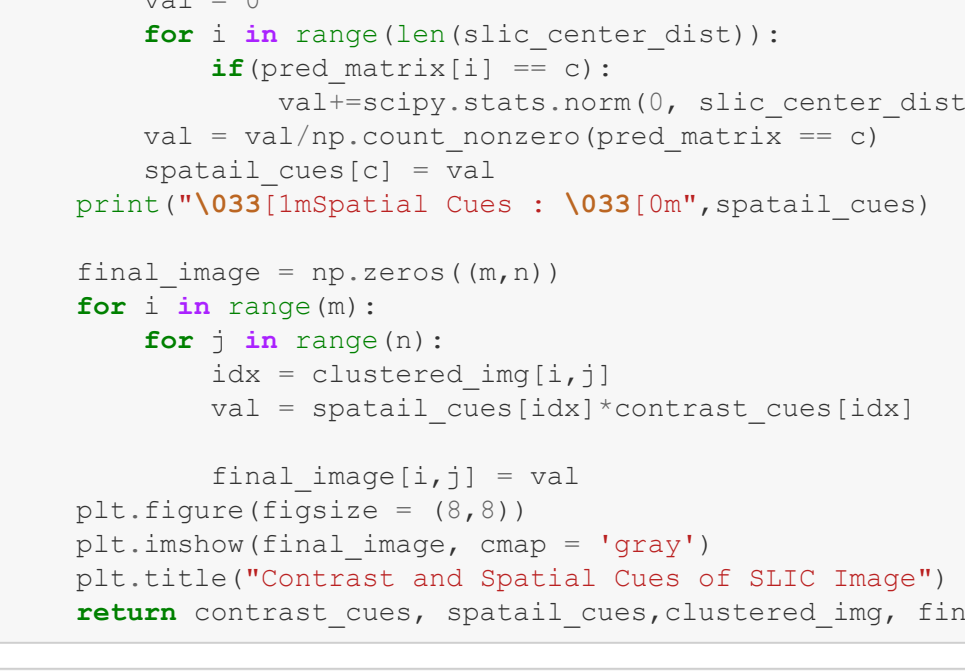
In [8]: `img = cv2.imread(img_path)
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
slic = cv2.ximgproc.createSuperpixelSLIC(img,region_size=15,ruler = 20.0)
slic.iterate(40) #Number of iterations, the greater the better
mask_slic = slic.getMask() #Get Mask, Super pixel edge Mask==1
label_slic = slic.getLabels() #Get superpixel tags
number_slic = slic.getNumberofSuperpixels() #Get the number of super pixels
cleaned_plus_img = cluster_centers[cleaned_img[:,0:3]]
img_slic = cv2.bitwise_and(img,img,mask = mask_inv_slic) #Draw the superpixel boundary on the original image
plt.figure(figsize = (8,8))
plt.imshow(img_slic)`

Out[8]: `<matplotlib.image.AxesImage at 0x24b5a18c588>`



In [9]: `def get_super_image(img, segments):
 n,n = segments.shape
 dict_ = {}
 centers = {}
 for i in range(n):
 for j in range(n):
 if(segments[i,j] not in dict_):
 dict_[segments[i,j]] = []
 centers[segments[i,j]] = []
 dict_[segments[i,j]].append(np.array([i,j]))
 centers[segments[i,j]].append(np.array([i,j]))
 else:
 dict_[segments[i,j]].append(np.array([i,j]))
 centers[segments[i,j]].append(np.array([i,j]))
 for key in list(dict_.keys()):
 dict_[key] = np.mean(np.array(dict_[key]), 0).astype(int)
 centers[key] = np.mean(np.array(centers[key]), 0).astype(int)
 slic_image = np.zeros((image.shape[0],image.shape[1],3))
 for i in range(n):
 for j in range(n):
 slic_image[i,j] = dict_[segments[i,j]]
 slic_image = slic_image.astype(int)
 plt.figure(figsize = (8,8))
 plt.imshow(slic_image)
 plt.title("SLIC Image")
 return slic_image, dict_, centers`

In [10]: `slic_image, dict_pixels, dict_centers = get_super_image(img,label_slic)`



Contrast Cue

Contrast cue represents the visual feature uniqueness on the single or multiple images. Contrast is one of the most widely used cues for measuring saliency in single image saliency detection algorithms, since the contrast operator simulates the human visual receptive fields. This rule is also valid in the case of cluster-based method for the multiple images. While the difference is that contrast cue on the cluster-level better represents the global correspondence relationship than the pixel/patch level.

The contrast cue $w^c(k)$ of cluster C^k is defined using its feature space as follows:

$$w^c(k) = \sum_{i=1, i \neq k}^K \left(\frac{n^k}{N} \|x^k - \mu^i\|_2 \right)$$

where a L2 norm is used to compute the distance on the feature space, n^k represents the pixel number of cluster C^k , and N denotes the pixel number of all images.

Spatial Cue

In human visual system, the regions near the image center draw more attention than the other regions. When the distance between the object and the image center increases, the attention gain is decreasing. This scenario is known as 'central bias rule' in single image saliency detection. We extend this concept to the cluster-based method, which measures a global spatial distribution of the cluster.

The spatial cue $w^s(k)$ of cluster C^k is defined as:

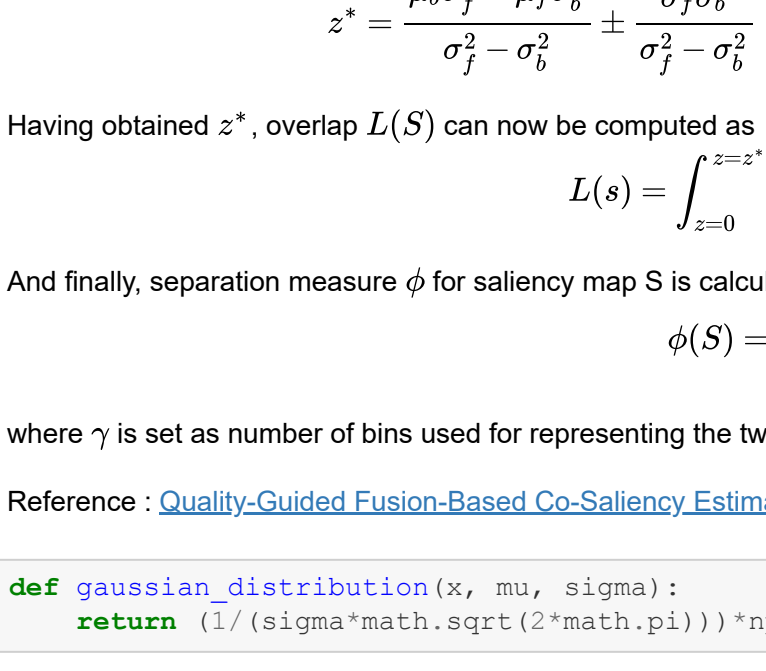
$$w^s(k) = \frac{1}{n^k} \sum_{i=1}^M \sum_{j=1}^{N_i} [N \|z^k - c^j\|^2 / (0, \sigma^2) \cdot \delta(b_j^k) - C^k]$$

where $\delta(\cdot)$ is the Kronecker delta function, c^j denotes the center of image I^j , and Gaussian kernel $N(\cdot)$ computes the Euclidean distance between pixel x_i^k and the image center c^j , the variance σ^2 is the normalized radius of images. And the normalization coefficient n^k is the pixel number of cluster C^k . Different from the single image model, our spatial cue w^s represents the location prior on the cluster-level, which is a global center bias on the multiple images. The same as the contrast cue, the spatial cue is also valid on both single and multiple images.

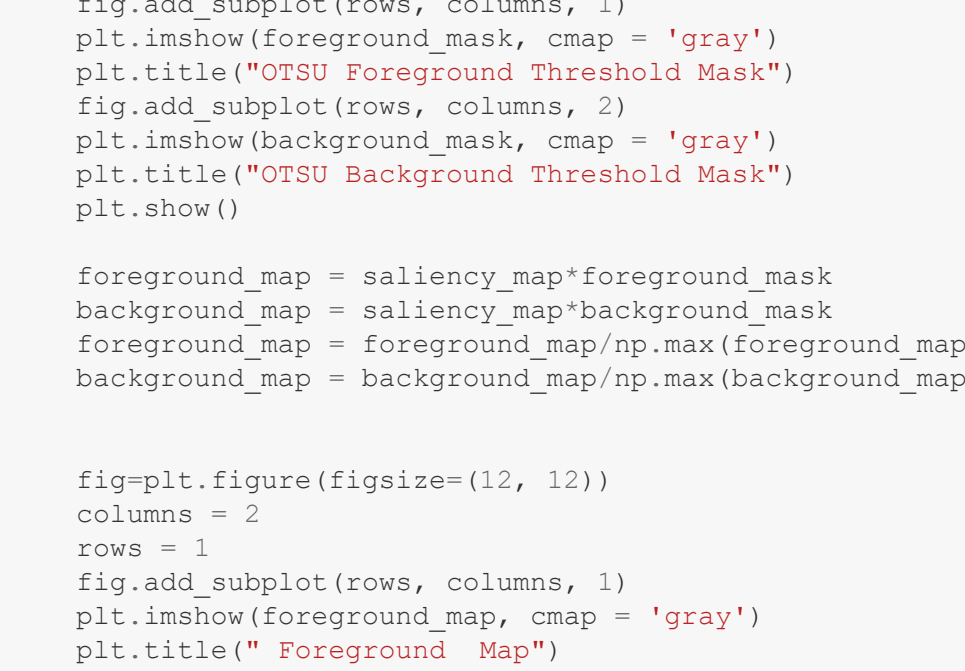
Reference : [Cluster-Based Co-Saliency Detection](#)

In [11]: `def contrast_spatial_cue(slic_image, dict_pixels, label_slic, dict_centers):
 n,n = label_slic.shape
 data = np.zeros((len(list(dict_pixels.keys())), 3))
 l = sorted(list(dict_pixels.keys()))
 for i in range(len(l)):
 data[i] = dict_pixels[l[i]]
 data = np.float32(data)/255
 s_scores = []
 k = 5
 s_max = 0
 for k,i in tqdm(range(10,61), position = 0, desc = "Finding best k"):
 criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 200, 0.2)
 _, labels, (centers) = cv2.kmeans(data, k, None, criteria, 200, cv2.RMEANS_RANDOM_CENTERS)
 ss = silhouette_score(data, np.squeeze(labels))
 s_scores.append(ss)
 if(ss > s_max):
 s_max = ss
 k = i
 plt.plot(l, [x for i in range(10,61)], s_scores)
 plt.xlabel('Value of k')
 plt.ylabel('silhouette score')
 plt.show()
 print("\033[1mbest Value of k obtained is : \033[0m",k)
 criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 200, 0.02)
 _, labels, (centers) = cv2.kmeans(data, k, None, criteria, 200, cv2.RMEANS_RANDOM_CENTERS)
 pred_matrix = np.squeeze(labels)
 cluster_dist = cdist(centers,centers, 'euclidean')
 clustered_img = pred_matrix[label_slic]
 n_i = ()
 for i in range(n):
 n_i[i] = np.count_nonzero(clustered_img == i)
 print("\033[1mNumber of Image Pixels per cluster : \033[0m",n_i)
 contrast_cues = {}
 for i in range(k):
 val = 0
 for j in range(k):
 val += (n_i[j]/(clustered_img.shape[0]*clustered_img.shape[1]))*(cluster_dist[i,j]
 contrast_cues[i] = val
 print("\033[1mContrast Cues : \033[0m",contrast_cues)
 center_data = np.zeros((len(list(dict_centers.keys())), 2))
 l = sorted(list(dict_centers.keys()))
 for i in range(len(l)):
 a,b = dict_centers[l[i]]
 center_data[i,0] = a
 center_data[i,1] = b
 slic_image_center = dict_centers[label_slic[n/2,n/2]].reshape((1,2))
 slic_image_center[0,0] = slic_image_center[0,1]
 slic_image_center[0,1] = slic_image_center[0,1]
 slic_center_dist = cdist(slic_image_center, 'euclidean')
 slic_center_dist_var = np.std(slic_center_dist)
 spatia_cues = {}
 for i in range(k):
 val = 0
 for i in range(len(slic_center_dist)):
 if(pred_matrix[i] == c):
 val = val/np.count_nonzero(pred_matrix == c)
 spatia_cues[c] = val
 print("\033[1mSpatial Cues : \033[0m",spatial_cues)
 final_image = np.zeros((m,n))
 for i in range(m):
 for j in range(n):
 idx = clustered_img[i,j]
 val = spatial_cues[idx]*contrast_cues[idx]
 final_image[i,j] = val
 plt.figure(figsize = (8,8))
 plt.imshow(final_image, cmap = 'gray')
 plt.imshow(dict_centers, cmap = 'gray')
 plt.title("Contrast and Spatial Cues of SLIC Image")
 return contrast_cues, spatial_cues,clustered_img, final_image`

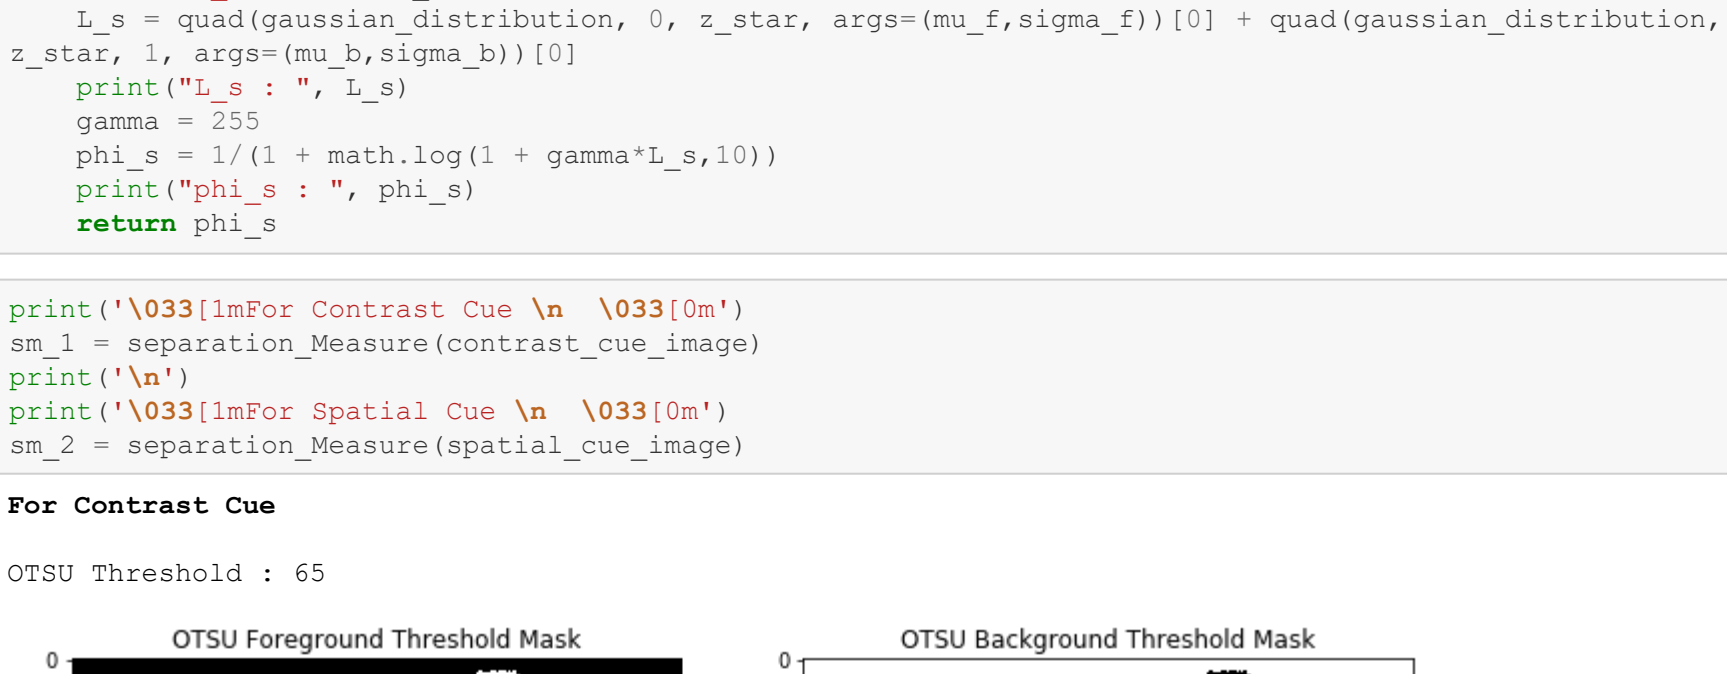
In [12]: `contrast_cues, spatial_cues,clustered_img, final_image = contrast_spatial_cue(slic_image, dict_pixels,
label_slic, dict_centers)
Finding best k: 100#
00:00.2321t/s
51/51 [00:21`



Best Value of k obtained is : 12
Number of Image Pixels per cluster : 10: 30850, 11: 18227, 12: 53983, 13: 31794, 14: 40286, 15: 23412, 16: 13144, 17: 25440, 18: 13962, 19: 23613, 20: 25370, 21: 103649
Contrast Cues : 10: 0.500177553295481, 11: 0.5745287569404574, 12: 0.4952804724575096, 13: 0.635933366589334, 14: 0.443165246076661, 15: 0.5724946495057592, 16: 0.5130627406302538, 17: 0.7372048309601246, 18: 0.46519157903491693, 19: 0.4373689002792746, 20: 0.4947235316349006, 21: 0.601163097676324, 22: 0.0014793389031498127, 23: 0.0005066217487978924, 24: 0.0004661763317128694, 25: 0.00046691656879523156, 26: 0.001820281254824394, 27: 0.0003240320224580659, 28: 0.0014447053618514187, 29: 0.0008471857270555645, 30: 0.0005815901929739707, 31: 0.000330164166210544



In [13]: `m,n = clustered_img.shape
contrast_cue_image = np.zeros((m,n))
spatial_cue_image = np.zeros((m,n))
for i in range(m):
 for j in range(n):
 contrast_cue_image[i,j] = contrast_cues[clustered_img[i,j]]
 spatial_cue_image[i,j] = spatial_cues[clustered_img[i,j]]
contrast_cue_image = (contrast_cue_image - np.min(contrast_cue_image))/(np.max(contrast_cue_image) - np.min(contrast_cue_image))
spatial_cue_image = (spatial_cue_image - np.min(spatial_cue_image))/(np.max(spatial_cue_image) - np.min(spatial_cue_image))
columns = 2
rows = 1
fig=plt.figure(figsize=(12, 12))
fig.add_subplot(rows, columns, 1)
plt.imshow(contrast_cue_image, cmap = 'gray')
plt.title("Contrast Cue Image")
fig.add_subplot(rows, columns, 2)
plt.imshow(spatial_cue_image, cmap = 'gray')
plt.title("Spatial Cue Image")
fig.tight_layout()
plt.show()`



Question 3

Implement the Otsu quality measures as discussed in Sec III.B.1 of the following paper to obtain quality scores for the two cues obtained in Q2. Use these quality scores as weights while performing the weighted sum of the two cues for getting the final saliency cue. [4 marks]

In [14]: `def otsu(img):
 min_cost = float('inf')
 threshold = 0
 for i in range(1,256):
 v0 = np.var(img[img < i], ddof = 1)
 w0 = len(img[img < i])
 v1 = np.var(img[img >= i], ddof = 1)
 w1 = len(img[img >= i])
 L_s = quad(gaussian_distribution, 0, z_star, args=(mu_f,sigma_f))[0] + quad(gaussian_distribution, z_star, 1, args=(mu_b,sigma_b))[0]
 cost = w0*v0 + w1*v1
 if(cost < min_cost):
 min_cost = cost
 threshold = i
 return threshold
def select_foreground(img):
 m,n = img.shape
 m0 = int(0.15*n)
 n0 = int(0.15*m)
 c0 = 0
 for i in range(m/2 - m0, m/2 + m0):
 for j in range(n/2 - n0, n/2 + n0):
 if(img[i,j] == 0):
 c0 = c0 + 1
 else:
 c1 = c1 + 1
 if(c0 < c1):
 return 0
 else:
 return 1
def separation_Measure(saliency_map):
 saliency_map = (saliency_map/np.max(saliency_map))*255
 thres = otsu(saliency_map,3,thresh)
 mask = copy.deepcopy(saliency_map)
 mask[mask < thres] = 0
 mask[mask >= thres] = 1
 fg = select_foreground(mask)
 if(fg == 1):
 foreground_mask = mask
 background_mask = 1 - foreground_mask
 else:
 foreground_mask = 1 - mask
 background_mask = 1 - foreground_mask
 fig=plt.figure(figsize=(12, 12))
 columns = 2
 rows = 1
 fig.add_subplot(rows, columns, 1)
 plt.imshow(foreground_mask, cmap = 'gray')
 plt.title("OTSU Foreground Threshold Mask")
 fig.add_subplot(rows, columns, 2)
 plt.imshow(background_mask, cmap = 'gray')
 plt.title("OTSU Background Threshold Mask")
 plt.show()
 foreground_map = saliency_map*foreground_mask
 background_map = saliency_map*background_mask
 foreground_map = foreground_map/np.max(foreground_map)
 background_map = background_map/np.max(background_map)
 fig=plt.figure(figsize=(12, 12))
 columns = 2
 rows = 1
 fig.add_subplot(rows, columns, 1)
 plt.imshow(foreground_map, cmap = 'gray')
 plt.title("Foregound Map")
 fig.add_subplot(rows, columns, 2)
 plt.imshow(background_map, cmap = 'gray')
 plt.title("Background Map")
 plt.show()
 mu_f = np.mean(foreground_map[foreground_map > 0])
 sigma_f = np.std(foreground_map[foreground_map > 0])
 mu_b = np.mean(background_map[background_map > 0])
 sigma_b = np.std(background_map[background_map > 0])
 print("foreground mean , mu_f : 0.6154494086655235
background mean , mu_b : 0.6024314003738889
foreground standard deviation , sigma_f : 0.31654428643427657
background standard deviation , sigma_b : 0.16304821932170072
z_star : 0.817037707922726
L_s : 0.5969327805330965
phi_s : 0.3139412464280025`

For Spatial Cue

OTSU Threshold : 65
OTSU Foreground Threshold Mask
OTSU Background Threshold Mask
Foregound Map
Background Map
foreground mean , mu_f : 0.6154494086655235
background mean , mu_b : 0.6024314003738889
foreground standard deviation , sigma_f : 0.31654428643427657
background standard deviation , sigma_b : 0.16304821932170072
z_star : 0.817037707922726
L_s : 0.5969327805330965
phi_s : 0.3139412464280025

In [18]: `def separation_Measure(saliency_map):
 saliency_map = (saliency_map/np.max(saliency_map))*255
 thres = otsu(saliency_map,3,thresh)
 mask = copy.deepcopy(saliency_map)
 mask[mask < thres] = 0
 mask[mask >= thres] = 1
 fg = select_foreground(mask)
 if(fg == 1):
 foreground_mask = mask
 background_mask = 1 - foreground_mask
 else:
 foreground_mask = 1 - mask
 background_mask = 1 - foreground_mask
 fig=plt.figure(figsize=(12, 12))
 columns = 2
 rows = 1
 fig.add_subplot(rows, columns, 1)
 plt.imshow(foreground_mask, cmap = 'gray')
 plt.title("Foregound Map")
 fig.add_subplot(rows, columns, 2)
 plt.imshow(background_mask, cmap = 'gray')
 plt.title("Background Map")
 plt.show()
 mu_f = np.mean(foreground_map[foreground_map > 0])
 sigma_f = np.std(foreground_map[foreground_map > 0])
 mu_b = np.mean(background_map[background_map > 0])
 sigma_b = np.std(background_map[background_map > 0])
 print("foreground mean , mu_f : 0.8310502841778272
background mean , mu_b : 0.4004934036498464
foreground standard deviation , sigma_f : 0.08659989638480675
background standard deviation , sigma_b : 0.2139826233903447
z_star : 0.6717269208054917
L_s : 0.0195047572430588
phi_s : 0.509866849735761`

In [18]: `final_saliency_image = np.zeros((m,n))
for i in range(m):
 for j in range(n):
 final_saliency_image[i,j] = sm_1*contrast_cues[clustered_img[i,j]] + sm_2*spatial_cues[clustered_img[i,j]]
plt.figure(figsize = (8,8))
plt.imshow(final_saliency_image, cmap = 'gray')
plt.title("Final Saliency Image")`

