

CSE 344 Computer Vision Assignment 2

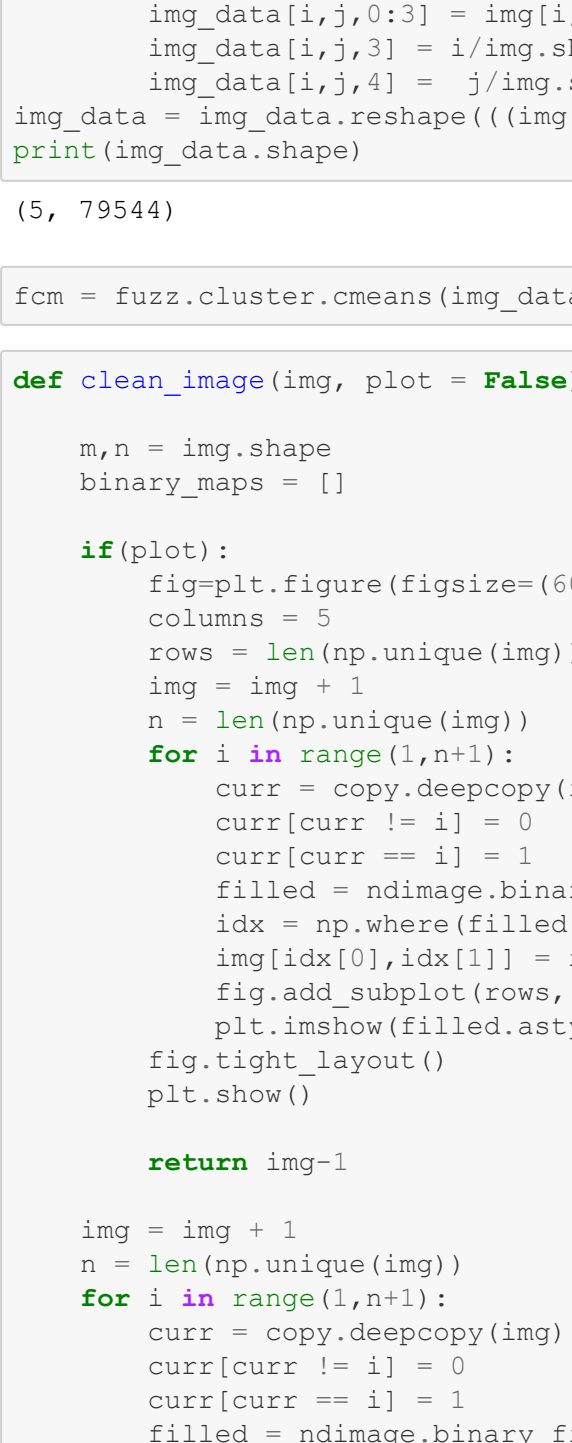
Name : Anka Sarkar
Roll Number : 2019222

```
In [2]: import numpy as np
import cv2
import matplotlib.pyplot as plt
from IPython.display import display
from PIL import Image
import skimage as skm
from skimage.metrics import silhouette_score
from skimage import io
from skimage.segmentation import slic
from skimage.segmentation import mark_boundaries
from skimage.util import img_as_float
from skimage import io
from scipy.spatial.distance import cdist
import scipy.stats
import copy
import math
from scipy.integrate import quad
from sklearn.cluster import KMeans
from scipy import ndimage
```

Question 1

Write a program to implement a region segmentation algorithm using the fuzzy c-means algorithm on normalized 'RGBxy' data of an image. Merge any isolated pixels (or very-small regions) to their surrounding regions. [3 marks]

```
In [45]: img = cv2.imread('Cats1.png')
fig=plt.figure(figsize=(6, 6))
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.show()
```



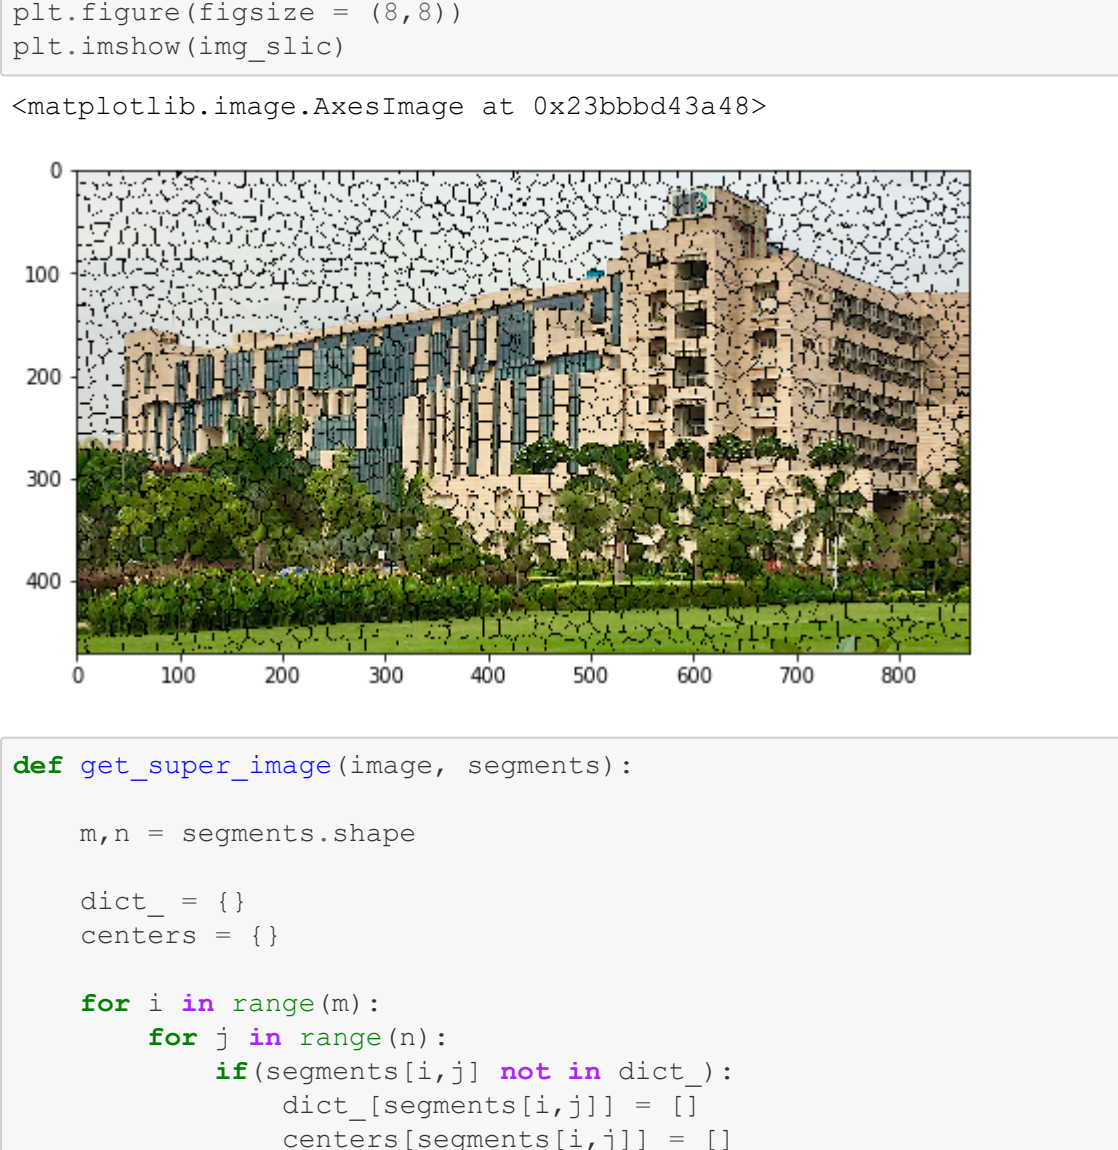
```
In [28]: img = np.array(img, dtype=np.float)
img_data = np.zeros((img.shape[0], img.shape[1], 5))
for i in range(img.shape[0]):
    for j in range(img.shape[1]):
        img_data[i,j,0] = img[i,j]/255
        img_data[i,j,1] = i/img.shape[0]
        img_data[i,j,4] = j/img.shape[1]
img_data = img_data.reshape((img.shape[0]*img.shape[1], 5)).T
print(img_data.shape)
```

```
In [29]: fcm = fuzz.cluster.cmeans(img_data, 25, 2, error=0.05, maxiter=1000, init=None)
```

```
In [30]: def clean_image(img, plot = False):
    m,n = img.shape
    binary_maps = []
    if(plot):
        fig=plt.figure(figsize=(60, 60))
        columns = 5
        rows = len(np.unique(img))/5
        img = img + 1
        n = len(np.unique(img))
        for i in range(1,n+1):
            curr = copy.deepcopy(img)
            cur[cure == i] = 0
            cur[cure == i] = 1
            filled = ndimage.binary_fill_holes(curr).astype(int)
            idx = np.where(filled == 1)
            img[idx[0],idx[1]] = i
            fig.add_subplot(rows, columns, i)
            plt.imshow(filled.astype(int))
            fig.tight_layout()
            plt.show()
        return img-1
    else:
        n = len(np.unique(img))
        for i in range(1,n+1):
            curr = copy.deepcopy(img)
            cur[cure == i] = 0
            cur[cure == i] = 1
            filled = ndimage.binary_fill_holes(curr).astype(int)
            idx = np.where(filled == 1)
            img[idx[0],idx[1]] = i
        return img-1
```

```
In [33]: cluster_centers = fcm[0]
prob_matrix = fcm[1]
pred_matrix = np.argmax(prob_matrix, axis = 0)
pred_matrix = cluster_centers[pred_matrix]
```

```
clustered_image = pred_matrix[:,0:3]
clustered_image = clustered_image.reshape((img.shape[0],img.shape[1], 3))*255
clustered_image = clustered_image.astype(int)
pred_matrix = np.argmax(prob_matrix, axis = 0)
cleaned_img = clean_image(np.squeeze(pred_matrix).reshape(img.shape[0],img.shape[1]))
cleaned_clus_img = Cluster_centers[cleaned_img[:, :,0:3]]
cleaned_clus_img = cleaned_clus_img.reshape((img.shape[0],img.shape[1], 3))*255
cleaned_clus_img = cleaned_clus_img.astype(int)
```

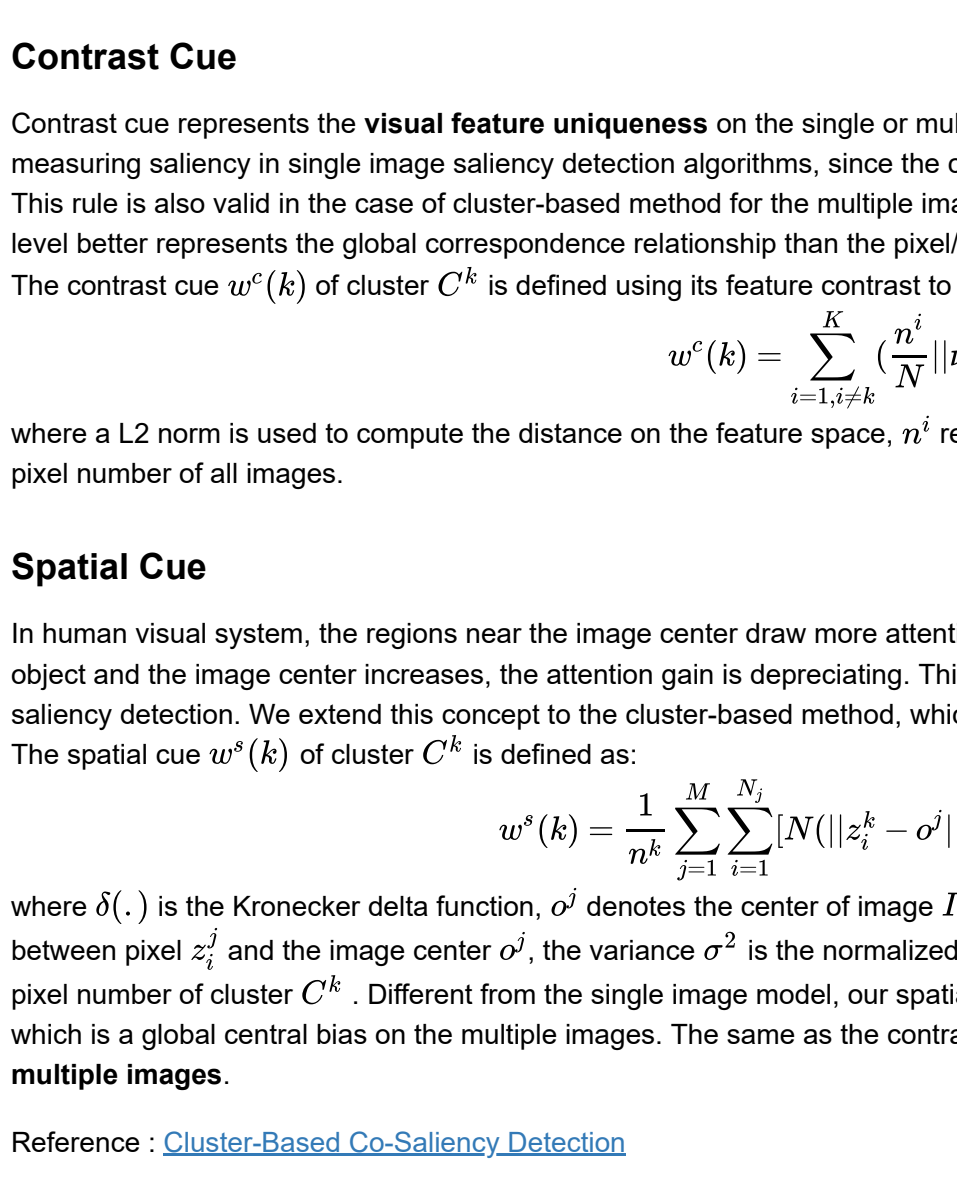


Question 2

Write a program to obtain the spatial and contrast cues using SLIC superpixels of an image instead of pixels. [3 marks]

```
In [41]: img_path = "1558014721_87jYwA_iiit_d.jpg"
image = cv2.imread(img_path)
plt.figure(figsize = (8,8))
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
print(image.shape)
```

```
Out [36]: <matplotlib.image.AxesImage at 0x23bbdb43a48>
```

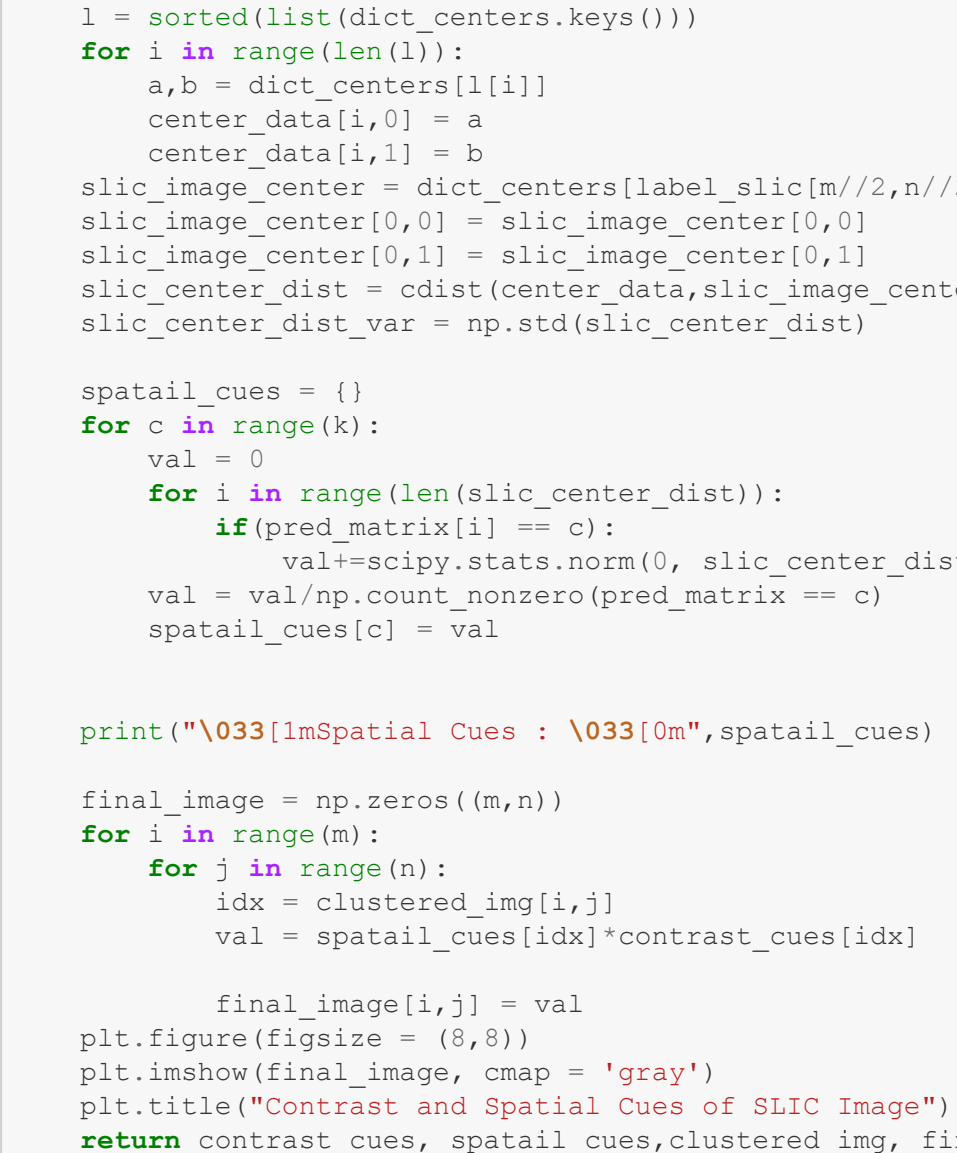


```
In [36]: img = cv2.imread(img_path)
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
#Initialize the slic item, the average size of super pixels is 20 (default is 10), and the smoothing fa
ctor is 20
slic = cv2.ximgproc.createSuperpixelSLIC(img,region_size=15,ruler = 20.0)
slic.iterate(40) #Number of iterations, the greater the better
mask_slic = slic.getLabelsContourMask() #Get Mask, Super pixel edge Mask==1
label_slic = slic.getLabels() #Get superpixel tags
mask_slic = slic.getNumberOFSuperpixels() #Get the number of super pixels
mask_inv_slic = cv2.bitwise_not(mask_slic)
label_slic = cv2.bitwise_and(img,mask mask_inv_slic) #Draw the superpixel boundary on the original
image
plt.figure(figsize = (8,8))
plt.imshow(label_slic)
```

```
Out [36]: <matplotlib.image.AxesImage at 0x23bbdb43a48>
```

```
In [39]: def get_super_image(image, segments):
    m,n = segments.shape
    dict_ = {}
    centers = []
    for i in range(m):
        for j in range(n):
            if(segments[i,j] not in dict_):
                dict_[segments[i,j]] = []
                centers[segments[i,j]] = []
            else:
                dict_[segments[i,j]].append(image[i,j])
                centers[segments[i,j]].append(np.array([i,j]))
    for key in list(dict_.keys()):
        dict_[key] = np.mean(np.array(dict_[key]), 0).astype(int)
        centers[key] = np.mean(np.array(centers[key]), 0).astype(int)
    slic_image = np.zeros((image.shape[0],image.shape[1],3))
    for i in range(m):
        for j in range(n):
            slic_image[i,j] = dict_[segments[i,j]]
    slic_image = slic_image.astype(int)
    plt.figure(figsize = (8,8))
    plt.imshow(slic_image)
    plt.title("SLIC Image")
    return slic_image, dict_, centers
```

```
In [40]: slic_image, dict_pixels, dict_centers = get_super_image(img,label_slic)
```



Contrast Cue

Contrast cue represents the visual feature uniqueness on the single or multiple images. Contrast is one of the most widely used cues for measuring saliency in single image saliency detection algorithms, since the contrast operator simulates the human visual receptive fields. This rule is also valid in the case of cluster-based method for the multiple images, while the difference is that contrast cue on the cluster-level better represents the cluster correspondence relationship than the pixel-level level.

The contrast cue $w^c(k)$ of cluster C^k is defined using its feature contrast to all other clusters:

$$w^c(k) = \frac{\sum_{i=1, i \neq k}^N (n^i)^{\frac{1}{2}} \| \mu^k - \mu^i \|^2}{\sum_{i=1, i \neq k}^N (n^i)^{\frac{1}{2}}}$$

where n is 2 norm is used to compute the distance on the feature space, n^i represents the pixel number of cluster C^i , and N denotes the pixel number of all images.

Spatial Cue

In human visual system, the regions near the image center draw more attention than the other regions. When the distance between the object and the image center increases, the attention gain is decreasing. This scenario is known as "central bias rule" in single image saliency detection. We extend this concept to the cluster-based method, which measures a global spatial distribution of the cluster. The spatial cue $w^s(k)$ of cluster C^k is defined as:

$$w^s(k) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N N(i,j) \left(\| z_i^k - z_j^k \|^2 + \sigma^2 \right) \delta(b_j^k) - C^k$$

where $\delta(\cdot)$ is the Kronecker delta function, z_i^k denotes the center of image of C^k , and Gaussian kernel $N(\cdot)$ computes the Euclidean distance between z_i^k and the image center z^0 , the variance σ^2 is the normalized radius of images. And the normalization coefficient n^k is the pixel number of cluster C^k . Different from the single image model, our spatial cue w^s represents the location prior on the cluster-level, which is a global central bias on the multiple images. The same as the contrast cue, the spatial cue is also valid on both single and multiple images.

Reference: [Cluster-Based Co-Saliency Detection](#)

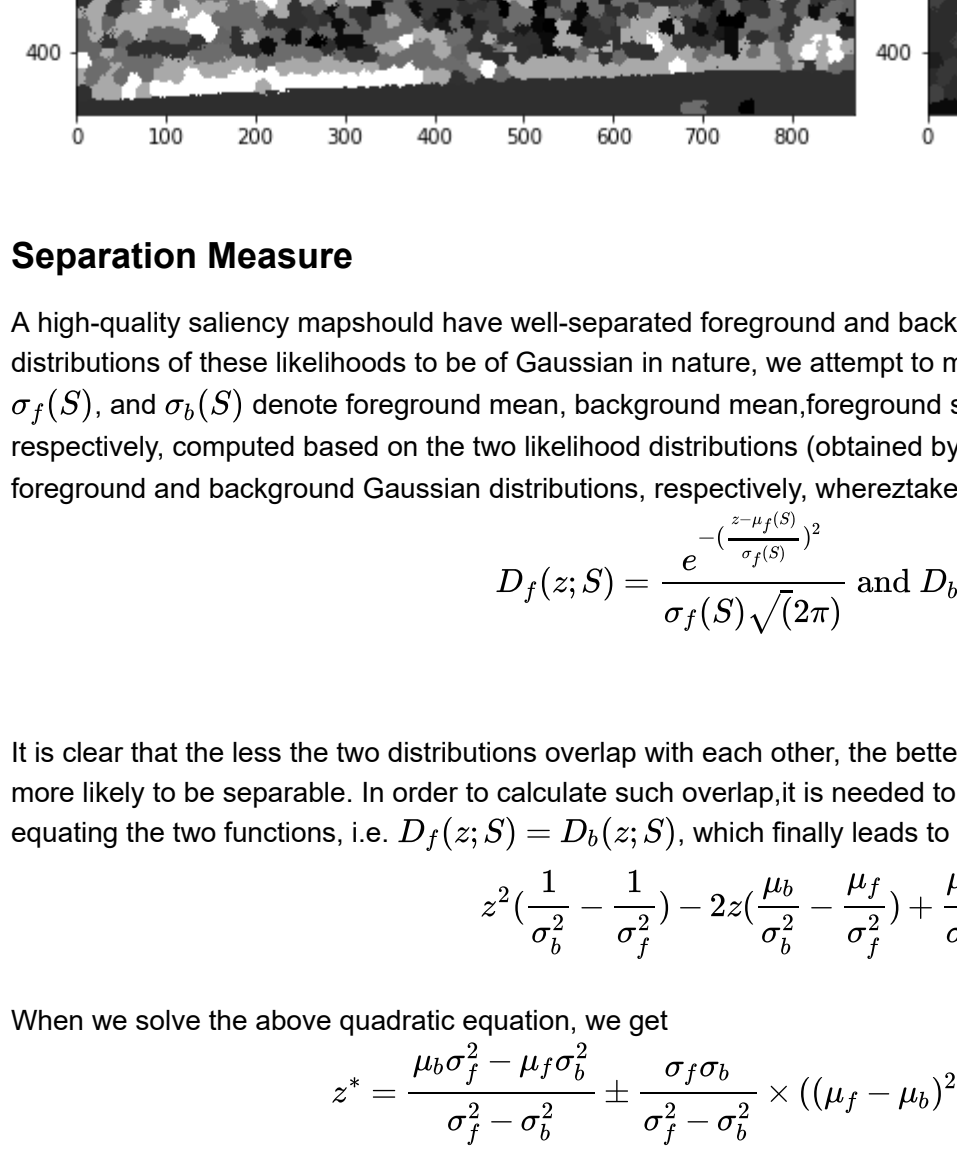
```
In [46]: def contrast_spatial_cue(slic_image, dict_pixels, label_slic, dict_centers):
    m,n = label_slic.shape
    data = np.zeros((len(list(dict_pixels.keys())), 3))
    l = sorted(list(dict_pixels.keys()))
    for i in range(len(l)):
        data[i] = dict_pixels[l[i]]
    data = np.float32(data)/255
    s_scores = []
    s_max = 0
    for k in range(1,61):
        criteria = cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 200, 0.2)
        labels, (centers) = cv2.kmeans(data, k, None, criteria, 200, cv2.KMEANS_RANDOM_CENTERS)
        ss = silhouette_score(data, np.squeeze(labels))
        s_scores.append(ss)
        if ss > s_max:
            s_max = ss
            k = k_i
    plt.plot([i for i in range(1,61)], s_scores)
    plt.xlabel("Value of k")
    plt.ylabel("Silhouette score")
    plt.show()
    print("\033[4mBest Value of k obtained is: %d\033[0m",k)
    criteria = cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 200, 0.2)
    labels, (centers) = cv2.kmeans(data, k, None, criteria, 200, cv2.KMEANS_RANDOM_CENTERS)
    pred_matrix = np.squeeze(labels)
    cluster_dist = cdist(centers,centers, 'euclidean')
    clustered_img = pred_matrix[label_slic]
    n_i = {}
    n_i[i] = np.count_nonzero(clustered_img == i)
    print("\033[4mNumber of Image Pixels per cluster : \033[0m",n_i)
    contrast_cues = {}
    for i in range(k):
        val = 0
        for j in range(k):
            if(i!=j):
                val += n_i[j]/(clustered_img.shape[0]*clustered_img.shape[1])*cluster_dist[i,j]
        contrast_cues[i] = val
    print("\033[4mContrast Cues : \033[0m",contrast_cues)
    center_data = np.zeros((len(list(dict_centers.keys())), 2))
    l = sorted(list(dict_centers.keys()))
    for i in range(len(l)):
        a,b = dict_centers[l[i]]
        center_data[i,0] = a
        center_data[i,1] = b
    slic_image_center = dict_centers[label_slic[m/2,n/2]].reshape((1,2))
    slic_image_center[0,0] = slic_image_center[0,1]
    slic_image_center[0,1] = slic_image_center[0,1]
    slic_center_dist = cdist(center_data,slic_image_center, 'euclidean')
    slic_center_dist_var = np.std(slic_center_dist)
```

```
In [18]: contrast_cues, spatial_cues,clustered_img, final_image = contrast_spatial_cue(slic_image, dict_pixels,
label_slic, dict_centers)
```

Finding best k: 100% 45/45 [00:13
x00:00, 3.51it/s]



Best Value of k obtained is : 17
Number of Image Pixels per cluster : (0: 15514, 1: 31062, 2: 26788, 3: 10953, 4: 19745, 5: 31740, 6: 7982, 7: 10369, 8: 10369, 9: 21198, 10: 18604, 11: 17659, 12: 5484, 13: 13022, 14: 13641, 15: 9541, 16: 17654)
Contrast Cues : (0: 0.49440749399293695, 1: 0.4722449659073994, 2: 0.5567239925962213, 3: 0.4578088282387095, 4: 0.753048836957364, 5: 0.5203890119232772, 6: 0.4431900045142752, 7: 0.4392563646432681, 8: 0.6014940160675252, 9: 0.002894904905817073, 10: 0.00191874350189343, 11: 0.5139412385154153, 12: 0.4429462354277046, 13: 0.485510440183346, 14: 0.4447145079227963, 15: 0.4753008723939895, 16: 0.644179361688751)
Spatial Cues : (0: 0.000491252134934001, 1: 0.0007969799423250584, 2: 0.000501836910440256, 3: 0.001932081519375824, 4: 0.00037037478049512, 5: 0.001969897463593991, 6: 0.000191445016253553, 7: 0.000388181845050597, 8: 0.74847339590507e+05, 9: 0.0004933776359412004, 10: 0.0007989662132898919, 11: 0.000505977964292548, 12: 0.00058034369954211, 13: 0.001773501919044747, 14: 0.0003235213528164026, 15: 0.0004283797930381034, 16: 0.000453364687668396)



Question 3

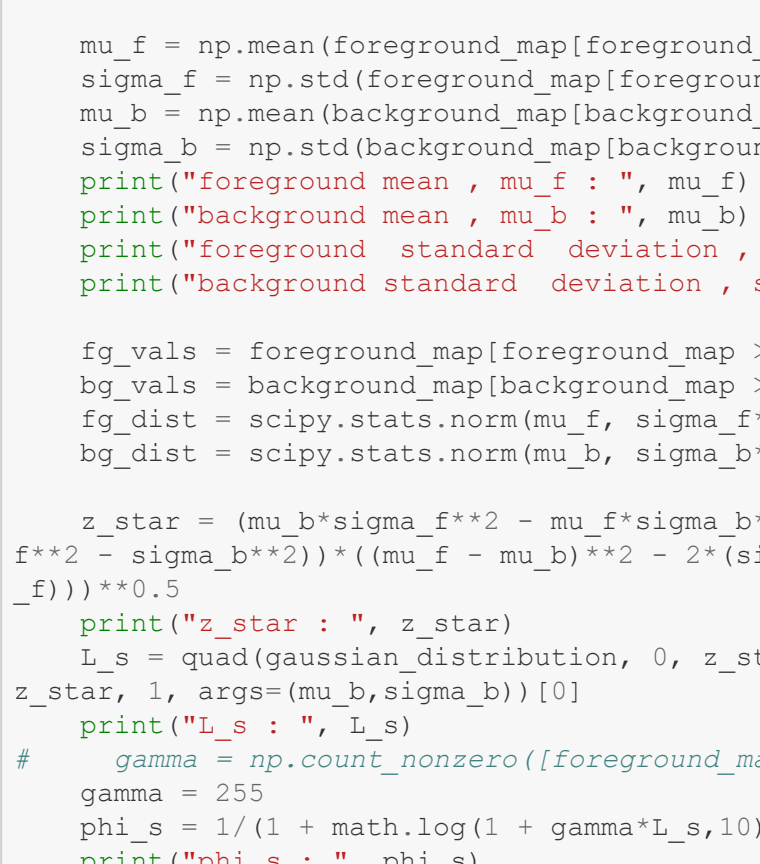
Implement the spatial-measure discussed in Sec III.B.1 of the following paper to obtain quality scores for the two cues obtained in Q2. Use these quality scores as weights while performing the weighted sum of the two cues for getting the final saliency cue. [4 marks]

```
In [47]: def otsu(img):
    min_cost = float('inf')
    threshold = 0
    for i in range(1,256):
        v0 = np.var(img[img < i], ddof = 1)
        v1 = len(img[img < i])
        w0 = np.var(img[img > i], ddof = 1)
        w1 = len(img[img > i])
        cost = w0*v0 + w1*v1
        if(min_cost > cost):
            min_cost = cost
            threshold = i
    return threshold
```

```
In [48]: def select_foreground(img):
    m,n = img.shape
    m0 = int(0.15*m)
    n0 = int(0.15*n)
    c0 = 0
    for i in range(m//2 - m0, m//2 + m0):
        for j in range(n//2 - n0, n//2 + n0):
            if(img[i,j] == 0):
                c0 = c0 + 1
            else:
                c1 = c1 + 1
    if(c0 > c1):
        return 0
    else:
        return 1
```

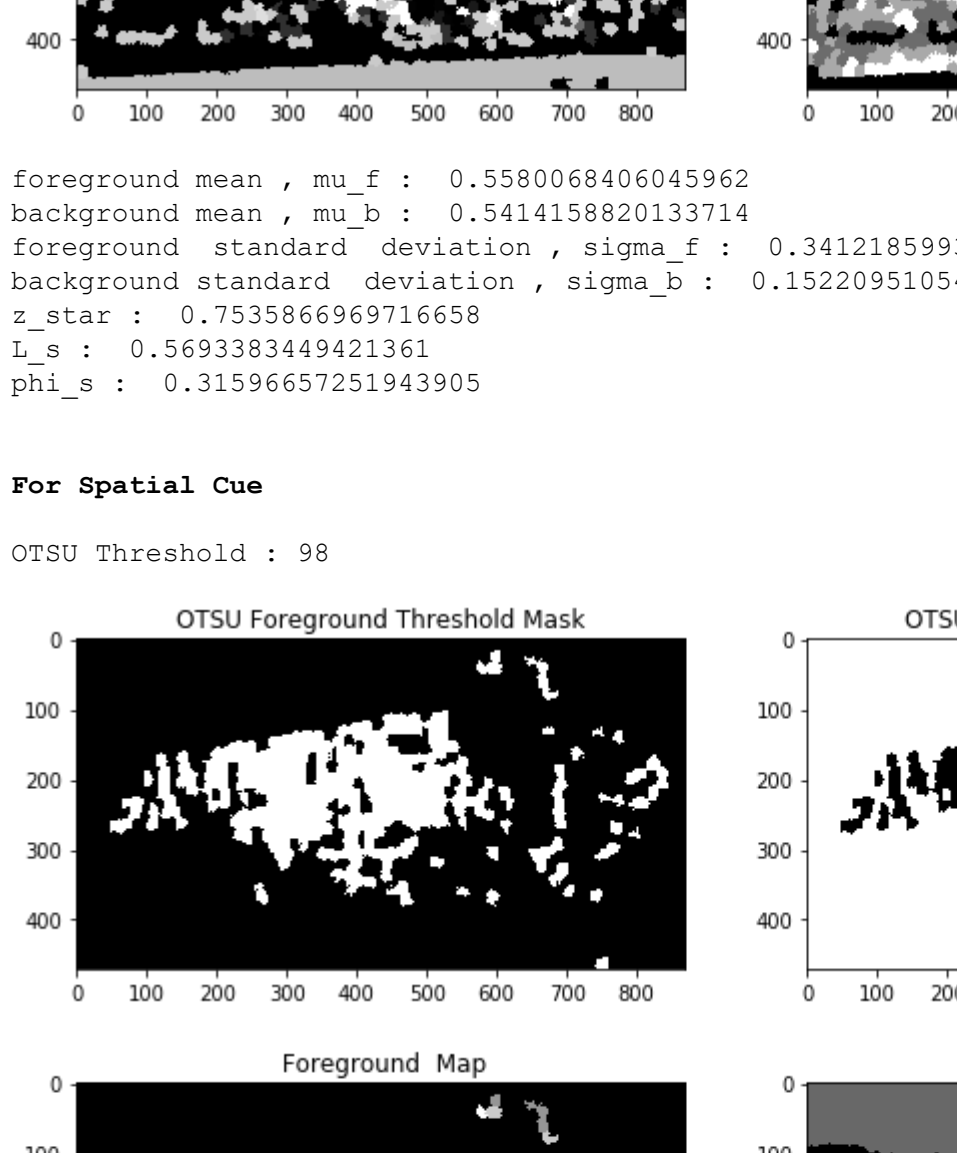
```
In [21]: contrast_cues, spatial_cues,clustered_img, final_image = contrast_spatial_cue(slic_image, dict_pixels,
label_slic, dict_centers)
```

Finding best k: 100% 45/45 [00:12
x00:00, 3.51it/s]



Best Value of k obtained is : 17
Number of Image Pixels per cluster : (0: 21746, 1: 12302, 2: 7951, 3: 10295, 4: 30620, 5: 30614, 6: 12047, 7: 81847, 8: 22193, 9: 36819, 10: 29734, 11: 26885, 12: 23393, 13: 18743, 14: 13641, 15: 2000, 16: 10506)
Contrast Cues : (0: 0.44900704399293695, 1: 0.44071208348229013, 2: 0.4416887414384442, 3: 0.4408430961078070, 4: 0.501092398974022, 5: 0.521970522610326, 6: 0.45601065910553085, 7: 0.41568377422702, 8: 0.574933066113366, 9: 0.00030104166210544, 10: 0.00044949532476071366, 11: 0.5827746304350001, 12: 0.000505977964292548, 13: 0.00058034369954211, 14: 0.001773501919044747, 15: 0.0003235213528164026, 16: 0.0004283797930381034, 17: 0.000453364687668396)

Spatial Cues : (0: 0.00037037478049512, 1: 0.000501836910440256, 2: 0.000491252134934001, 3: 0.001932081519375824, 4: 0.00037037478049512, 5: 0.001969897463593991, 6: 0.000191445016253553, 7: 0.000388181845050597, 8: 0.74847339590507e+05, 9: 0.0004933776359412004, 10: 0.0007989662132898919, 11: 0.000505977964292548, 12: 0.00058034369954211, 13: 0.001773501919044747, 14: 0.0003235213528164026, 15: 0.0004283797930381034, 16: 0.000453364687668396)



```
In [54]: m,n = clustered_img.shape
contrast_cue_image = np.zeros((m,n))
spatial_cue_image = np.zeros((m,n))
for i in range(m):
    for j in range(n):
        contrast_cue_image[i,j] = contrast_cues[clustered_img[i,j]]
        spatial_cue_image[i,j] = spatial_cues[clustered_img[i,j]]
for i in range(m):
    for j in range(n):
        spatial_cue_image[i,j] = spatial_cues[clustered_img[i,j]]
spatial_cue_image = (spatial_cue_image - np.min(spatial_cue_image))/(np.max(spatial_cue_image) - np.min(spatial_cue_image))
rows = 2
columns = 2
fig=plt.figure(figsize=(12, 12))
fig.add_subplot(rows, columns, 1)
plt.imshow(contrast_cue_image, cmap = 'gray')
plt.title("Contrast Cue Image")
fig.add_subplot(rows, columns, 2)
plt.imshow(spatial_cue_image, cmap = 'gray')
plt.title("Spatial Cue Image")
fig.tight_layout()
plt.show()
```



Separation Measure

A high-quality saliency map should have well-separated foreground and background likelihoods like a ground-truth binary mask. Assuming distributions of these likelihoods to be of Gaussian in nature, we attempt to measure the separation between the two $L_2(z; S)$, $\mu_0(S)$, $\sigma_1(S)$ and $\sigma_0(S)$ denote foreground mean, background mean foreground standard deviation, and background standard deviation respectively, computed based on the two likelihood distributions (obtained by Otsu thresholding). Let us denote $D_f(z; S)$ and $D_b(z; S)$ as foreground and background Gaussian distributions, respectively, wheretzates saliency value ranging between 0 and 1. Specifically,

$$D_f(z; S) = \frac{1}{\sigma_f(S)\sqrt{2\pi}} e^{-\frac{(z-\mu_f)^2}{2\sigma_f^2}} \text{ and } D_b(z; S) = \frac{1}{\sigma_b(S)\sqrt{2\pi}} e^{-\frac{(z-\mu_b)^2}{2\sigma_b^2}}$$

It is clear that the less the two distributions overlap with each other, the better the saliency map is, i.e., the foreground and background are more likely to be separable. In $D_f(z; S) = D_b(z; S)$, which finally leads to

$$z^2 \left(\frac{1}{\sigma_b^2} - \frac{1}{\sigma_f^2} \right) - 2z \left(\frac{\mu_f}{\sigma_b^2} - \frac{\mu_b}{\sigma_f^2} \right) + \frac{\mu_f^2}{\sigma_b^2} - \frac{\mu_b^2}{\sigma_f^2} + 2 \log \left(\frac{\sigma_b}{\sigma_f} \right) = 0$$

When we solve the above quadratic equation, we get

$$z^* = \frac{\mu_b \sigma_f^2 - \mu_f \sigma_b^2}{\sigma_f^2 - \sigma_b^2} \pm \frac{\sigma_f \sigma_b}{\sigma_f^2 - \sigma_b^2} \times \left((\mu_f - \mu_b)^2 - 2(\sigma_f^2 - \sigma_b^2) (\log(\sigma_b) - \log(\sigma_f)) \right)^{\frac{1}{2}}$$

Having obtained z^* , overlap $L(S)$ can now be computed as

$$L(s) = \int_{z=0}^{z^*} D_f(z; S) dz + \int_{z^*}^{z=1} D_b(z; S) dz$$

And finally, separation measure ϕ for saliency map S is calculated as

$$\phi(S) = \frac{1}{1 + \log_{10}(1 + \gamma L(S))}$$

where γ is set as number of bins used for representing the two distributions.

Reference: [Gaussian-Guided Fusion-Based Co-Saliency Estimation for Image Co-Segmentation and Colocalization](#)

```
In [56]: def gaussian_distribution(mu, sigma):
    return 1/(sigma*math.sqrt(2*math.pi))*np.exp(-(x/sigma - mu/sigma)**2)
```

```
In [57]: def separation_Measure(saliency_map):
    saliency_map = (saliency_map - np.min(saliency_map))/255
    thresh = otsu(saliency_map.astype(int))
    print("OTSU Threshold: ", thresh)
    mask = copy.deepcopy(saliency_map)
    mask[mask >= thresh] = 1
    fg = select_foreground(mask)
    if(fg == 1):
        foreground_mask = mask
        background_mask = 1 - foreground_mask
    else:
        foreground_mask = 1 - mask
        background_mask = 1 - foreground_mask
```

```
fig=plt.figure(figsize=(12, 12))
columns = 2
rows = 1
fig.add_subplot(rows, columns, 1)
plt.imshow(foreground_mask, cmap = 'gray')
plt.title("OTSU Foreground Threshold Mask")
fig.add_subplot(rows, columns, 2)
plt.imshow(background_mask, cmap = 'gray')
plt.title("OTSU Background Threshold Mask")
plt.show()
```

```
foreground_map = saliency_map*foreground_mask
background_map = saliency_map*background_mask
foreground_map = foreground_map/np.max(foreground_map)
background_map = background_map/np.max(background_map)
```

```
fig=plt.figure(figsize=(12, 12))
columns = 2
rows = 1
fig.add_subplot(rows, columns, 1)
plt.imshow(foreground_map, cmap = 'gray')
plt.title("Foreground Map")
fig.add_subplot(rows, columns, 2)
plt.imshow(background_map, cmap = 'gray')
plt.title("Background Map")
plt.show()
```

```
mu_f = np.mean(foreground_map[foreground_map > 0])
sigma_f = np.std(foreground_map[foreground_map > 0])
mu_b = np.mean(background_map[background_map > 0])
sigma_b = np.std(background_map[background_map > 0])
print("foreground mean , mu_f : ", mu_f)
print("background mean , mu_b : ", mu_b)
print("foreground standard deviation , sigma_f : ", sigma_f)
print("background standard deviation , sigma_b : ", sigma_b)
```

```
fg_vals = foreground_map[foreground_map > 0].flatten()
bg_vals = background_map[background_map > 0].flatten()
fg_dist = scipy.stats.norm(mu_f, sigma_f**2).pdf(fg_vals)
bg_dist = scipy.stats.norm(mu_b, sigma_b**2).pdf(bg_vals)
```

```
z_star = (mu_b*sigma_f**2 - mu_f*sigma_b**2)/(sigma_f**2 - sigma_b**2) + (sigma_f*sigma_b/(sigma_f**2 - sigma_b**2))*((mu_f - mu_b)**2 - 2*(sigma_f**2 - sigma_b**2)*(math.log(sigma_b) - math.log(sigma_f)))
L_s = quad(gaussian_distribution, 0, z_star, args=(mu_f,sigma_f))[0]
print("L_s : ", L_s)
gamma = np.count_nonzero(foreground_map > 0) + np.count_nonzero(background_map > 0)
phi_s = 1/(1 + math.log(1 + gamma*L_s))
print("phi_s : ", phi_s)
return phi_s
```

```
In [58]: final_saliency_image = np.zeros((m,n))
sm_1 = separation_Measure(contrast_cue_image)
print("\n")
print("\033[4mInfo Spatial Cue An \033[0m",sm_1)
sm_2 = separation_Measure(spatial_cue_image)
```

For Contrast Cue

OTSU Threshold : 63


```
foreground mean , mu_f : 0.580069406045962
background mean , mu_b : 0.541415982013714
foreground standard deviation , sigma_f : 0.341218559367307
background standard deviation , sigma_b : 0.152209510541837
L_s : 0.569338449421361
phi_s : 0.31596657251943905
```

For Spatial Cue

OTSU Threshold : 98


```
foreground mean , mu_f : 0.8808924741487429
background mean , mu_b : 0.50224654020149
foreground standard deviation , sigma_f : 0.1425002644191153
background standard deviation , sigma_b : 0.2308980359587217
z_star : 0.6902573963272953
L_s : 0.110222601050113
phi_s : 0.4058428943995639
```

```
In [59]: final_saliency_image = np.zeros((m,n))
for i in range(m):
    for j in range(n):
        final_saliency_image[i,j] = sm_1*contrast_cues[clustered_img[i,j]] + sm_2*spatial_cues[clustered_img[i,j]]
final_saliency_image = final_saliency_image - np.min(final_saliency_image)/(np.max(final_saliency_image) - np.min(final_saliency_image))
plt.figure(figsize = (8,8))
plt.imshow(final_saliency_image, cmap = 'gray')
plt.title("Final Saliency Image")
```

Out [59]: Text(0.5, 1.0, 'Final Saliency Image')

